

Cascading

www.cascading.org

info@cascading.org

Design Goals

- [Make large processing jobs more transparent
- [Reusable processing components independent of resources
- [Incremental “data” builds
- [Simplify testing of processes
- [Scriptable from higher level languages (Groovy, JRuby, Jython, etc)

Cascading Introduction

Tuple Streams

— Tuple

— A set of ordered data ["John", "Doe", 39]

— Value Stream

— Just tuples

— Group Stream

— Tuples groups by a key

Value Stream

[V1,V2,...,Vn

[V1,V2,...,Vn

[V1,V2,...,Vn

[V1,V2,...,Vn

[V1,V2,...,Vn

[V1,V2,...,Vn

[V1,V2,...,Vn

Group Stream

[K1,K2,...,Kn

[V1,V2,...,Vn

[V1,V2,...,Vn

[V1,V2,...,Vn

[K1,K2,...,Kn

[V1,V2,...,Vn

[V1,V2,...,Vn

[V1,V2,...,Vn

[V1,V2,...,Vn

Tuple Streams

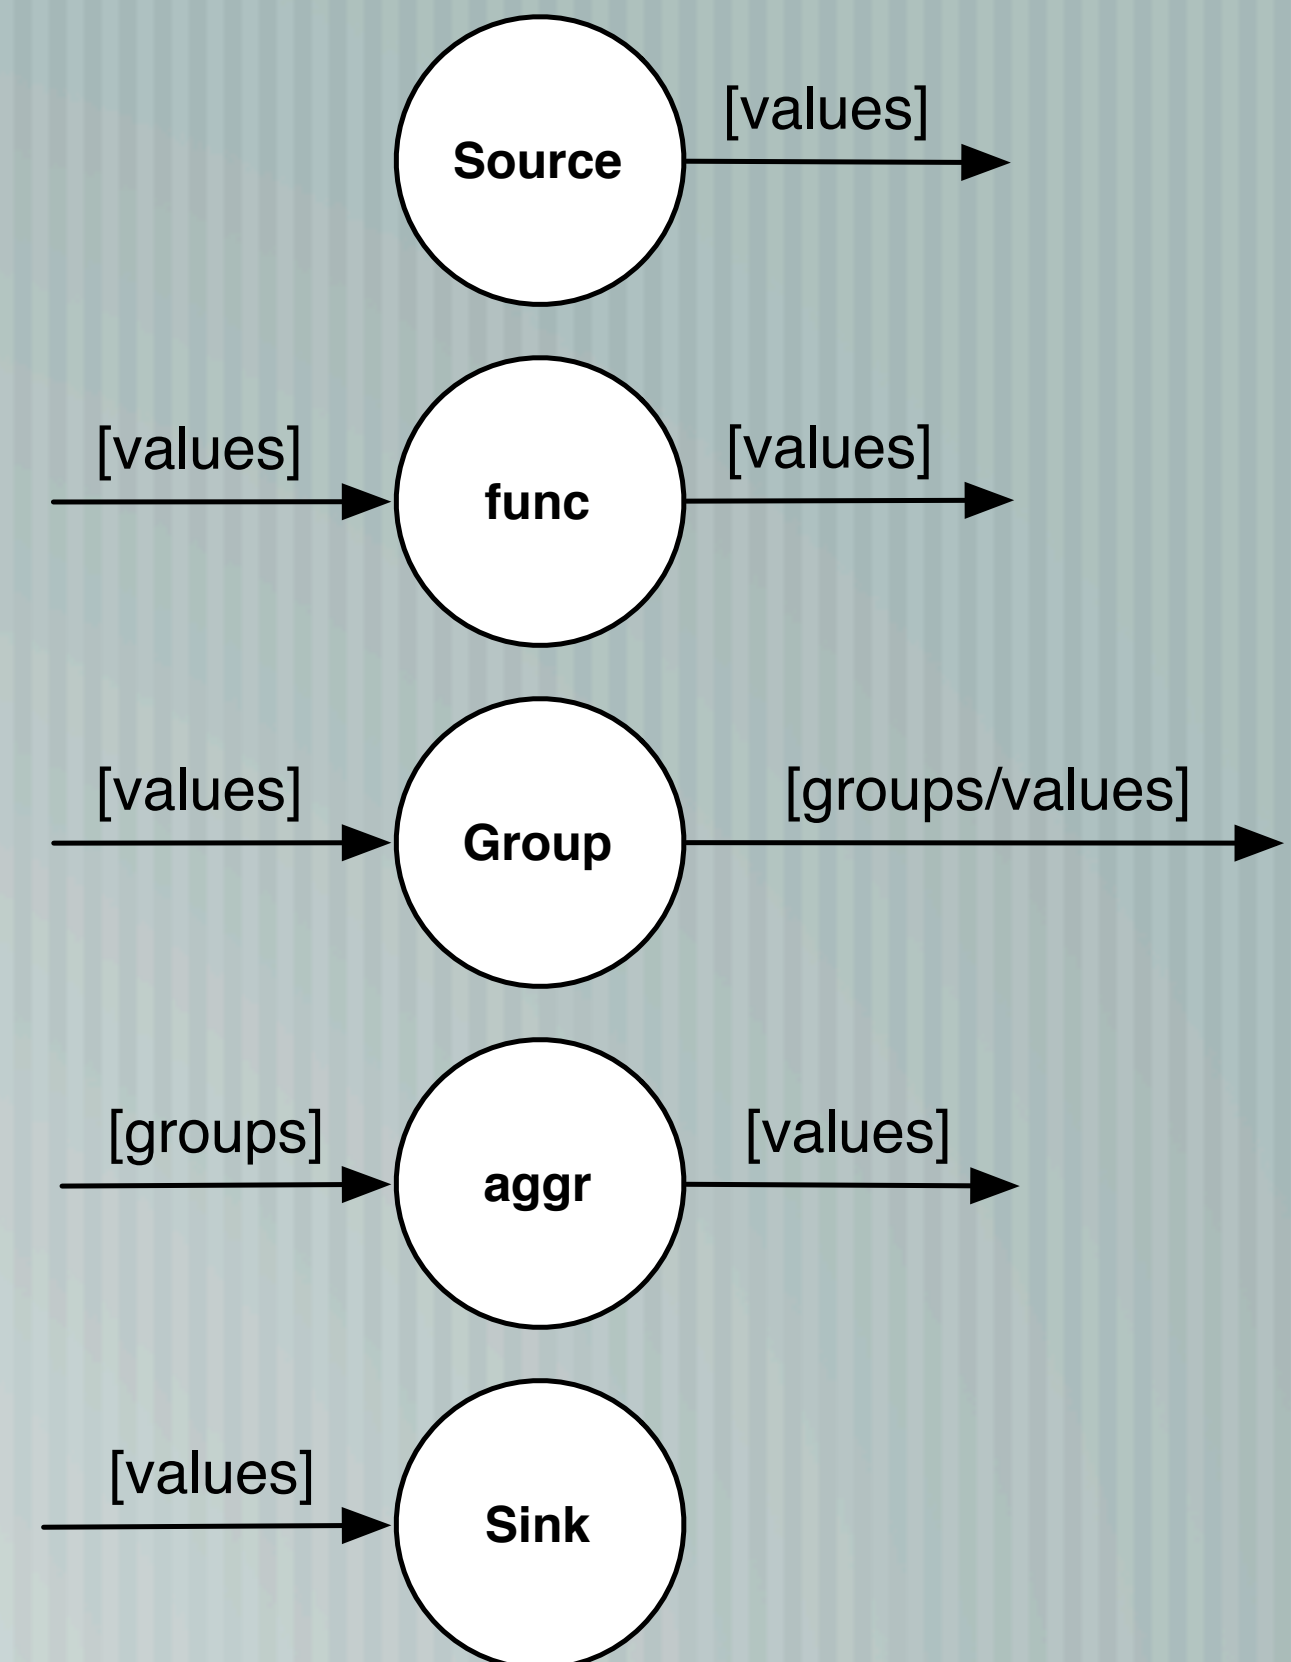

- Scalar functions and filters

- Apply to value and group streams


- Aggregate functions

- Apply to group stream

- Functions can be chained



Stream Processing



Pipe Assemblies


- A chain of scalar functions, groupings, aggregate functions
- Reusable, independent of data source/sink

Flows

- Assemblies plus sources and sinks



Cascades

- A collection of Flows





Processing Patterns


Chain




Splits



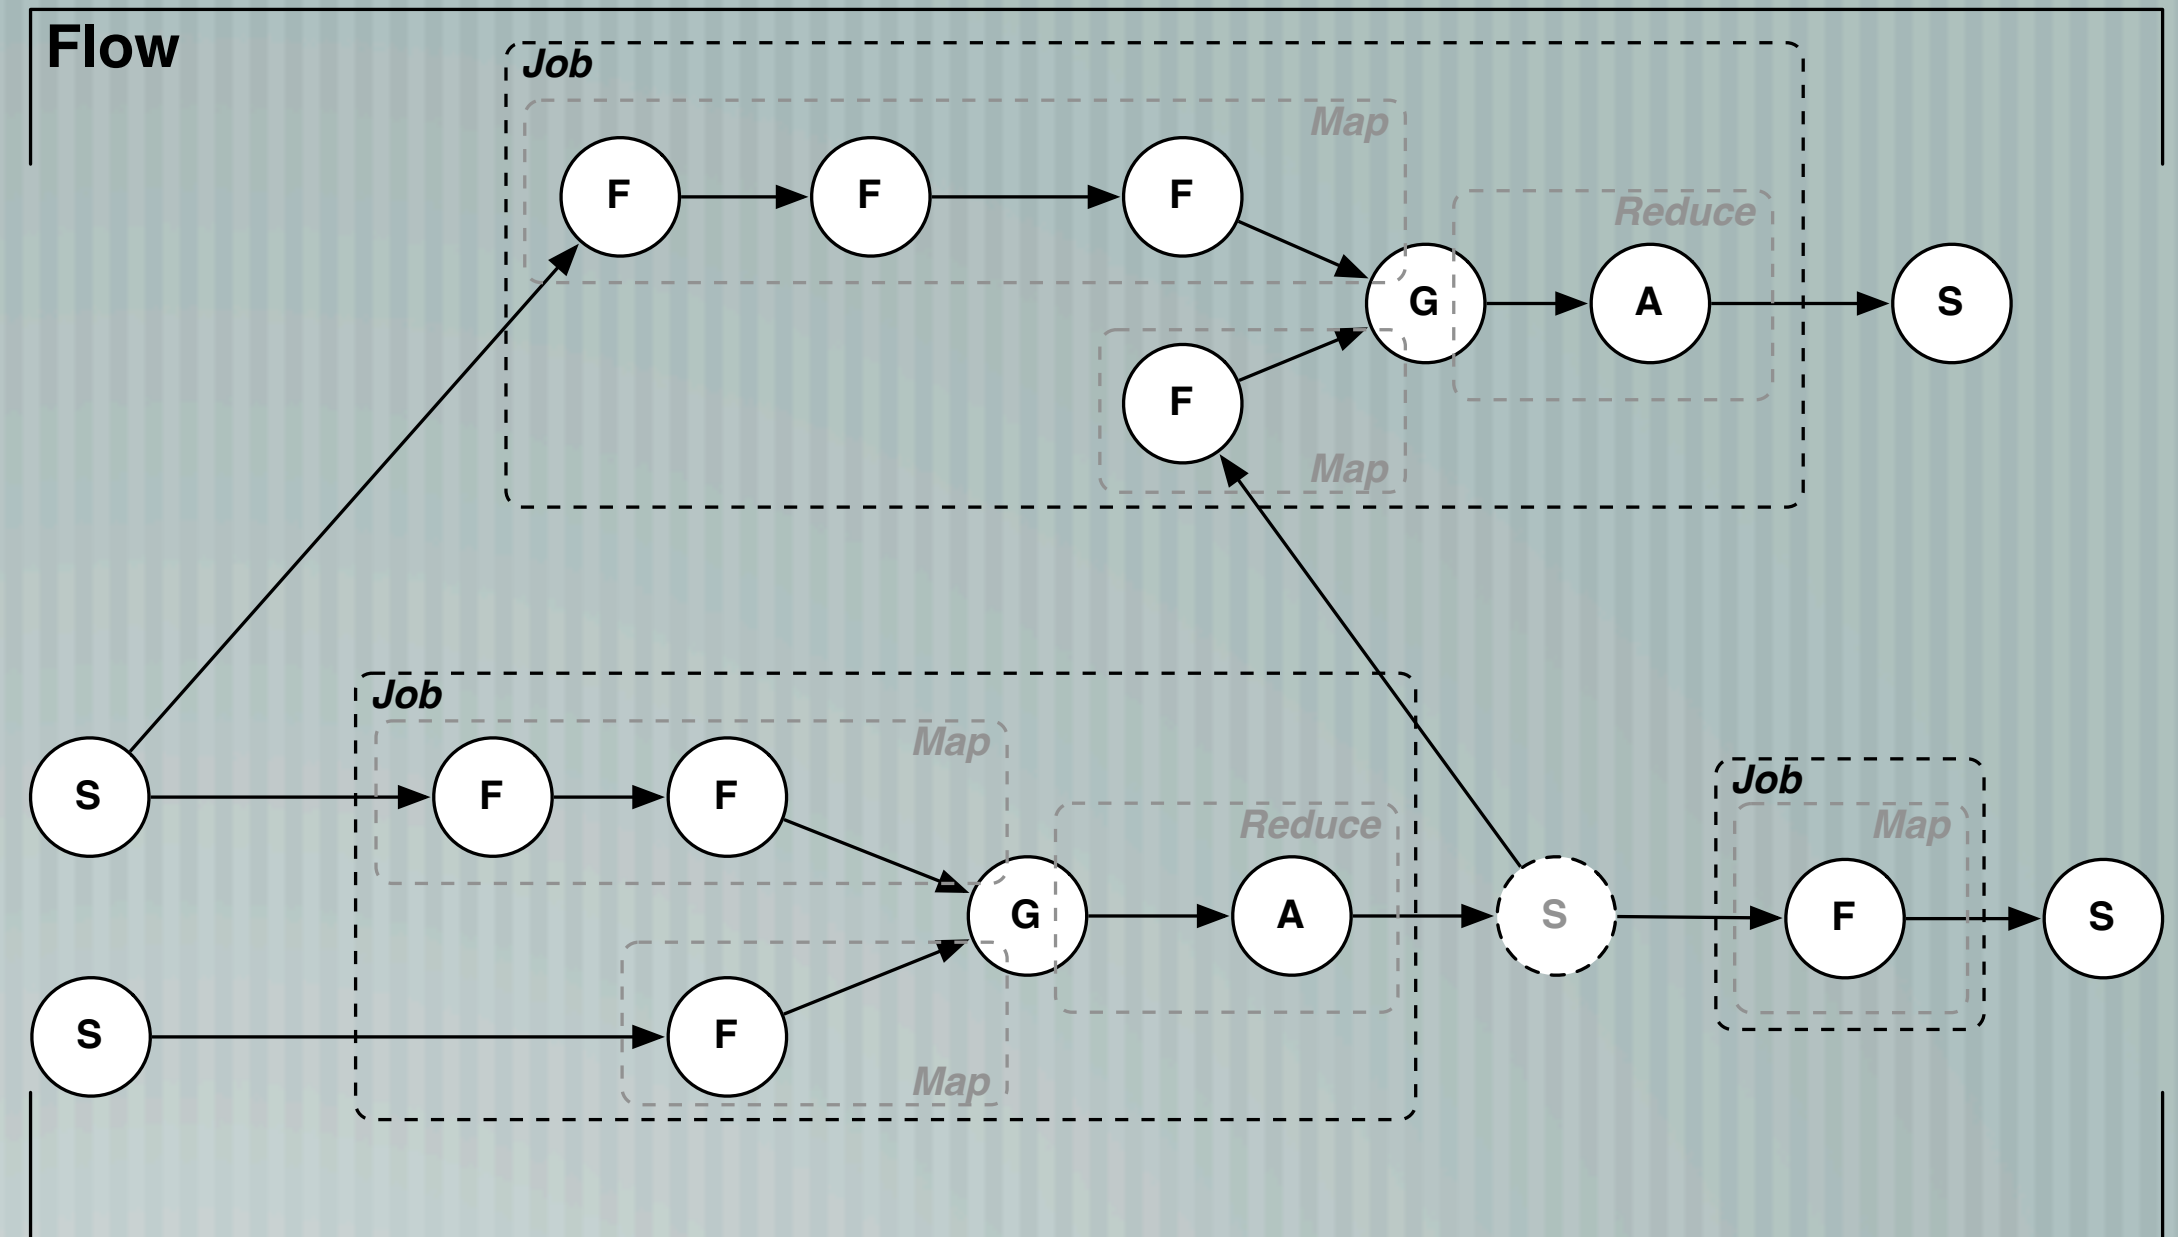

Joins



Cross



MapReduce Planner



Flows are logical 'units of work'

Flows 'compiled' into MR Jobs

Intermediate files are created (and destroyed) to join Jobs

Topological Scheduler

- [Flows walk MapReduce Jobs in dependency order
- [Cascades walk Flows in dependency order
- [Independent Jobs and Flows are scheduled to run concurrently
- [Listeners can react to element events (notify completion or failures)
- [Only stale data-sets are rebuilt (configurable)

Scripting - Groovy

```
Flow flow = builder.flow("wordcount")
{
    source(input, scheme: text()) // input is filename of raw text document

    tokenize(/[.,]*\s+/) // output new tuple for each split, result replaces stream by default
    group() // group on stream
    count() // count values in group, creates 'count' field by default
    group(["count"], reverse: true) // group/sort on 'count', reverse the sort order

    sink(output)
}

flow.complete() // execute, block till completed
```

System Integration

- [FileSystems (unique to Cascading)
 - Raw file S3 reading/writing (MD5)
 - Raw file HTTP reading (MD5)
 - Zip files
 - Can bypass native Hadoop 'collectors'
- [Event notification via listeners (XMPP/SQS/Zookeeper notifications)
- [Groovy scripting for easier local shell/file operations (wget, scp, etc)

Cascading API & Internals

Core Concepts

- [Taps and Schemes
- [Tuples and Fields
- [Pipes and PipeAssemblies
- [Each and Every Operators
- [Groups
- [Flows, FlowSteps, and FlowConnectors
- [Cascades, and CascadeConnectors, optional

Taps and Schemes


- [Taps, abstract out where and how a data resources is accessed
 - hdfs, http, local, S3, etc
- [Taps, used as Tuple (data) stream sinks, sources, or both
- [Schemes, define what a resource is made of
 - text lines, SequenceFile, CSV, etc

Tuples and Fields

- [Tuples are the 'records', read from Tap sources, written to Tap sinks
- [Fields are the 'column names', sourced from Schemes
- [Tuple class, an ordered collection of Comparable values
 - ("a string", 1.0, new SomeComparableWritable())
- [Fields class, a list of field names, absolute or relative positions
 - ("total", 3, -1) // fields 'total', 4th position, last position



Pipes and PipeAssemblies

- [Tuple streams pass through Pipes to be processed
- [Pipes, apply functions, filters, and aggregators to the Tuple stream
- [Pipe instances are chained together into assemblies
- [Reusable assemblies are subclasses of class PipeAssembly



Group Class and Subclasses

- Group, subclass of Pipe, groups the Tuple stream on given fields
- GroupBy and CoGroup subclass Group
- GroupBy groups and sorts
- CoGroup performs joins




Each and Every Classes

—— [Each, subclass of Pipe, applies Functions and Filters to each Tuple instance

—— $(a,b,c) \rightarrow \text{Each}(\text{func}()) \rightarrow (a,b,c,d)$

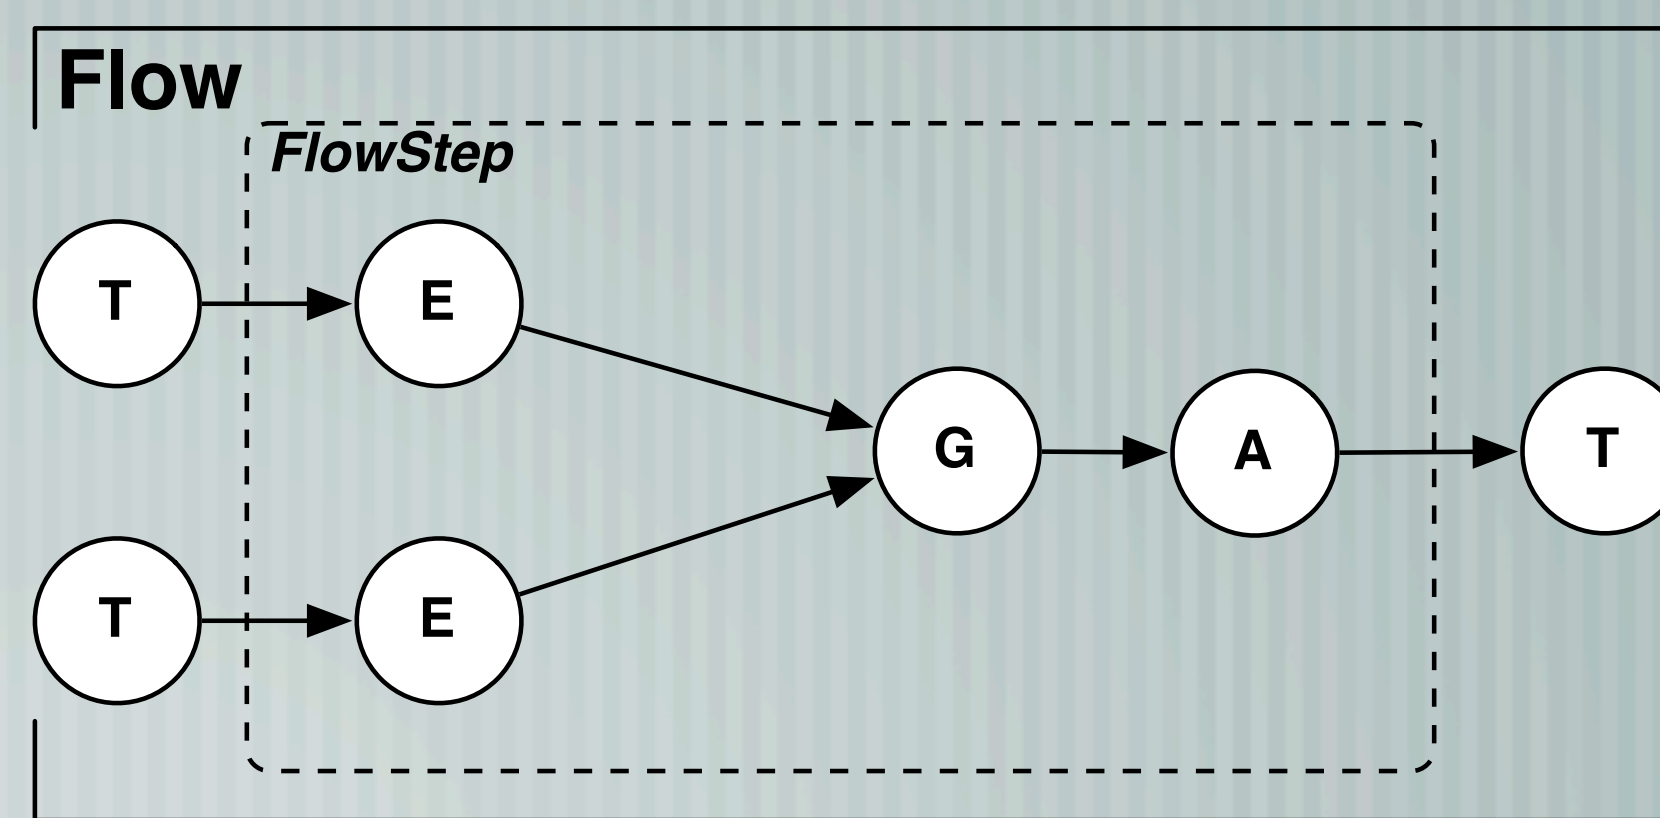
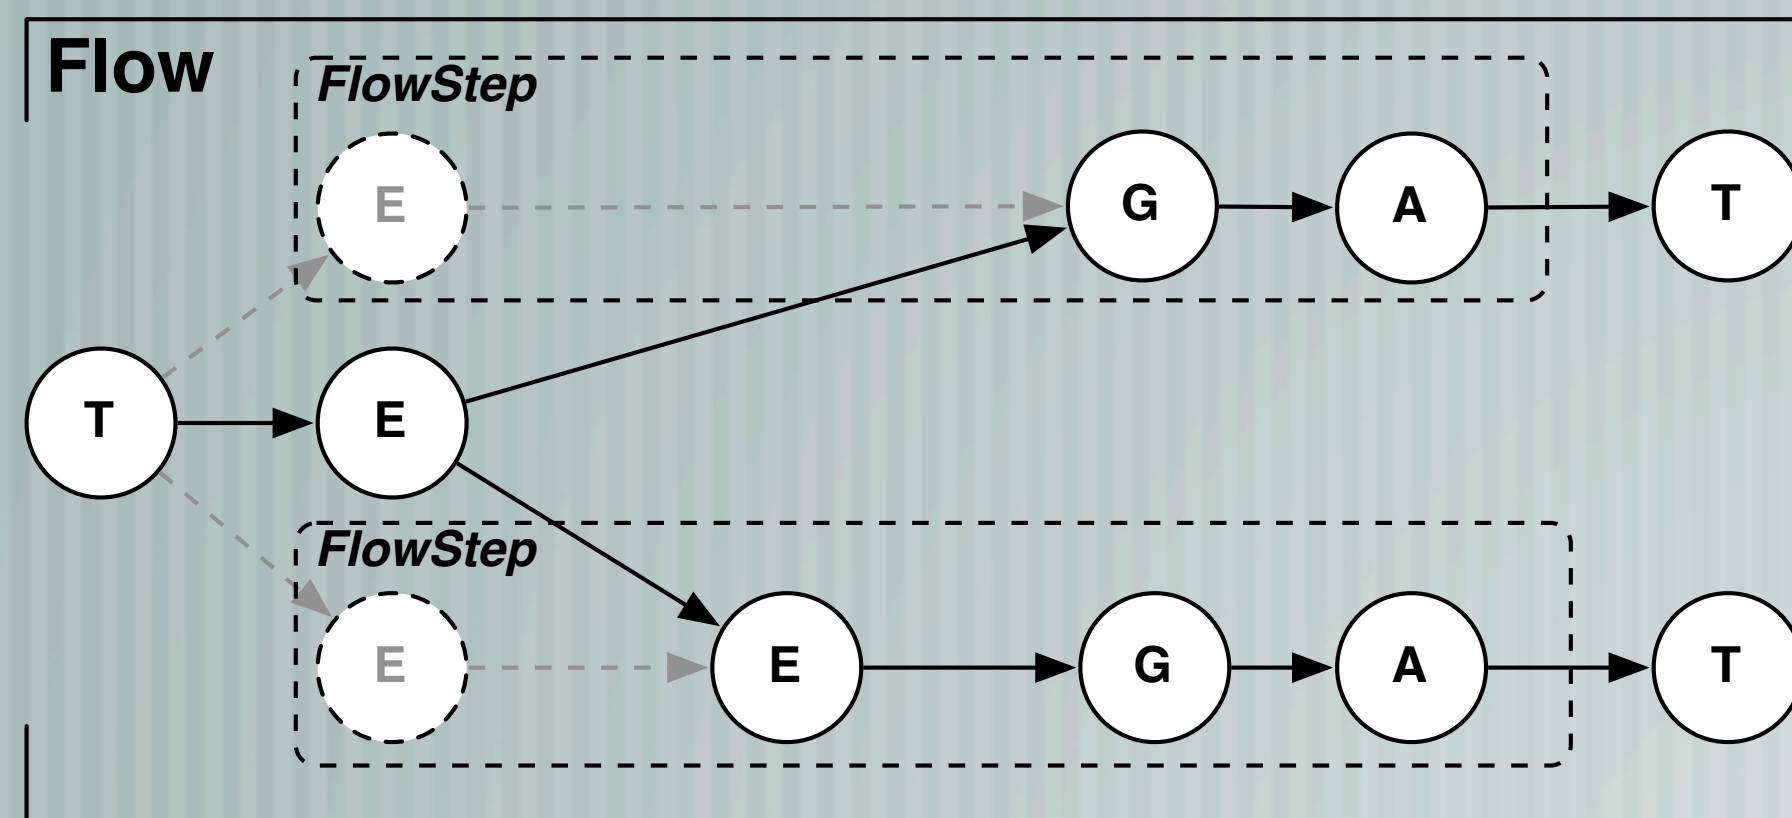
—— [Every, subclass of Pipe, applies Aggregators to every Tuple group

—— $(a: b,c) \rightarrow \text{Every}(\text{agg}()) \rightarrow (a,d: b,c)$




Flows and FlowConnectors

- Flows encapsulate assemblies and sink and source Taps
- FlowConnectors connect assemblies and Taps into Flows



FlowSteps and FlowConnectors

- Internally, FlowConnectors 'compile' assemblies into FlowSteps
- FlowSteps are MapReduce jobs, which are executed in Topo order
- Temporary files are created to link FlowSteps




Cascades and CascadeConnectors

— [Are optional

— [Cascades bind Flows together via shared Taps

— [CascadeConnectors connect Flows

— [Flows are executed in Topo order



Syntax

- [Each(previous, argSelector, function/filter, resultSelector)
- [Every(previous, argSelector, aggregator, resultSelector)
- [GroupBy(previous, groupSelector, sortSelector)
- [CoGroup(joinN, joiner, declaredFields)
- [Function(numArgs, declaredFields,)
- [Filter (numArgs, ...)
- [Aggregator(numArgs, declaredFields, ...)